


# Adaptive Interview Logic

After every answer, one score decides two things: **should the next question be a follow-up?** and **should it be easier or harder?** Both decisions are made by the app's own rules, then Gemini writes the question.

## Flowchart: from answer score to next question



# Score → decision at a glance

| ANSWER SCORE                       | NEXT QUESTION                   | NEXT DIFFICULTY                    |
|------------------------------------|---------------------------------|------------------------------------|
| Below 4                            | Follow-up on the same topic     | One level lower (Easy stays Easy)  |
| 4 to 6.99                          | Follow-up on the same topic     | Baseline                           |
| 7 to 8.5                           | New topic by priority           | Baseline                           |
| Above 8.5                          | New topic by priority           | One level higher (Hard stays Hard) |
| Any score, right after a follow-up | New topic (no second follow-up) | Back to baseline                   |

## Design notes

- **Two independent decisions, one score.** Follow-up-ness and difficulty are computed separately, so a very weak answer gets a follow-up that is also easier.
- **One step, then reset.** Difficulty moves by at most one level and returns to the candidate's chosen baseline on the next question, so the interview never drifts too far.
- **One follow-up per topic.** A weak answer earns a single targeted follow-up built from the earlier feedback and improvement tip, which avoids repetitive drilling.
- **Topic coverage by priority.** Fresh topics from the resume-based plan are used first (high → medium → low), then topics asked only once, until the 10-question limit or the candidate ends early.

Implemented in InterviewPage.jsx (selectNextTopic, selectNextDifficulty) · React, Node.js, Express, Gemini, Supabase PostgreSQL · [Live demo](#) · [GitHub](#)